


Intro to CNN





Left to right: Yann LeCun, Geoff Hinton, Yoshua Bengio, Andrew Ng at NIPS 2014 (from Andrew Ng's Facebook page).

1998
LeCun et al.




of transistors

 10^6

of pixels used in training

10⁷ NIST

2012
Krizhevsky et al.




of transistors GPUs

 10^9

of pixels used in training

 10^{14} 



Hubel & Wiesel, 1959

Simple cells:
Response to light
orientation

Complex cells:
Response to light
orientation and movement

Hypercomplex cells:
Response to movement with
end point




No response



Response
(end point)

Stimulus



Electrical signal
from brain




- 1.조각을 본다
- 2.각 조각이 조합된 패턴을 본다
- 3.점점 더 복잡한 조합의 패턴을 본다.
- 4.반응하는 여러 패턴의 조합을 가지고 이미지를 인식한다.

Learning of object parts

Examples of learned object parts from object categories



Faces



Cars




Elephants



Chairs





“필터를 CNN이 스스로 학습한다.”

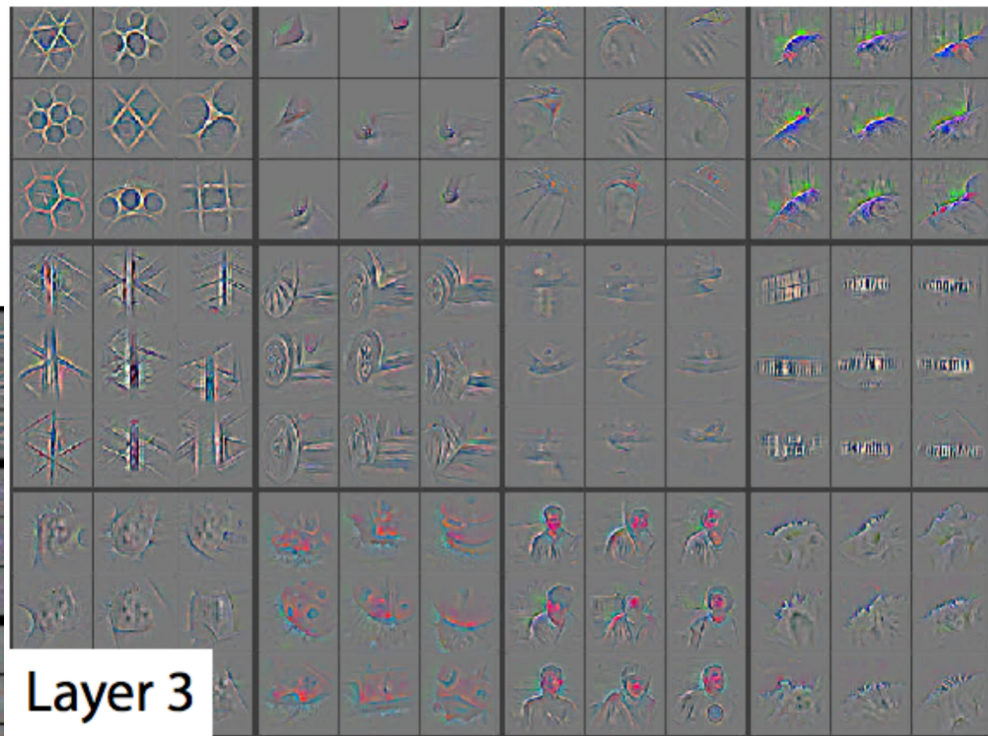




Layer 1




Layer 2



Layer 3

Convolution Layer

32x32x3 image




5x5x3 filter



Convolve the filter with the image
i.e. “slide over the image spatially,
computing dot products”

Convolution Layer

32x32x3 image




Filters always extend the full depth of the input volume

5x5x3 filter




Convolve the filter with the image
i.e. "slide over the image spatially,
computing dot products"

Convolution Layer




Convolution Layer

consider a second, **green** filter



For example, if we had 6 5x5 filters, we'll get 6 separate activation maps:



We stack these up to get a “new image” of size 28x28x6!

Convolutional Layer

Filter

for Conv Layer 01 filter - shape=(5, 5, 3)

W1 = tf.Variable(tf.random_normal([5, 5, 3, 32], stddev=0.01))


for Conv Layer 02 filter - shape=(5, 5, 32)

W2 = tf.Variable(tf.random_normal([5, 5, 32, 64], stddev=0.01))



A closer look at spatial dimensions:

7



7x7 input (spatially)
assume 3x3 filter

7

A closer look at spatial dimensions:

7


7

7x7 input (spatially)
assume 3x3 filter

=> 5x5 output

A closer look at spatial dimensions:

7




7

7x7 input (spatially)
assume 3x3 filter
applied **with stride 2**

A closer look at spatial dimensions:

7



7

7x7 input (spatially)
assume 3x3 filter
applied **with stride 2**
=> 3x3 output!

In practice: Common to zero pad the border

0	0	0	0	0	0			
0								
0								
0								
0								

e.g. input 7x7

3x3 filter, applied with **stride 1**

pad with 1 pixel border => what is the output?

7x7 output!

Convolutional Layer

Padding

Convolution Layer 01 -> (?, 28, 28, 32)

```
L1 = tf.nn.conv2d(X_img, W1, strides=[1, 1, 1, 1], padding='SAME')
```

```
L1 = tf.nn.relu(L1)
```


Pooling Layer 01 -> (?, 14, 14, 32)

```
L1 = tf.nn.max_pool(L1, ksize=[1, 2, 2, 1], strides=[1, 2, 2, 1], padding='SAME')
```


참고: input과 같은 size 28x28의 결과를 얻기 위해서
padding 1이 내부적으로 적용 되었음.



Preview: ConvNet is a sequence of Convolution Layers, interspersed with activation functions



Preview: ConvNet is a sequence of Convolutional Layers, interspersed with activation functions



Convolutional Layer

Activation function

Convolution Layer 01 -> (?, 28, 28, 32)

```
L1 = tf.nn.conv2d(X_img, W1, strides=[1, 1, 1, 1], padding='SAME')
```

```
L1 = tf.nn.relu(L1)
```


Pooling Layer 01 -> (?, 14, 14, 32)

```
L1 = tf.nn.max_pool(L1, ksize=[1, 2, 2, 1], strides=[1, 2, 2, 1], padding='SAME')
```




Pooling layer

- makes the representations smaller and more manageable
- operates over each activation map independently:



MAX POOLING



Pooling Layer

Sub sampling - Max Pooling

Convolution Layer 01 -> (?, 28, 28, 32)

```
L1 = tf.nn.conv2d(X_img, W1, strides=[1, 1, 1, 1], padding='SAME')
```

```
L1 = tf.nn.relu(L1)
```


Pooling Layer 01 -> (?, 14, 14, 32)

```
L1 = tf.nn.max_pool(L1, ksize=[1, 2, 2, 1], strides=[1, 2, 2, 1], padding='SAME')
```



Fully Connected Layer

32x32x3 image -> stretch to 3072 x 1



Fully Connected Layer

Flatten

flatten

```
L_FC = tf.reshape(L2_pool, [-1, 7 * 7 * 64], name="L_flat")
```

logits = $W_{fc} * L_{FC} + \text{bias}$

```
logits = tf.add(tf.matmul(L_FC, W_fc), b, name='logits')
```



Cost & Optimizer

Softmax

Cost(loss) function & Optimizer

```
cost = tf.reduce_mean(  
    tf.nn.softmax_cross_entropy_with_logits(  
        logits=logits, labels=Y))  
optimizer = tf.train  
    .AdamOptimizer(learning_rate=learning_rate)  
    .minimize(cost)
```



